

EE379K: Assignment #2

Due: Mar. 1 2024, 11:59pm TX

Introduction

The overall goal of this homework is to implement a local feature matching algorithm. You can discuss in small groups, but turn in individual solutions and indicate collaborators. Turn in assignments by **Friday, Mar. 1**. Submit to Canvas a .zip file containing (1) a questions.pdf with answers for Part 1 questions, (2) a writeup.pdf for Part 2 code; and (3) a code/ directory containing your code without any large result files or subfolders. The pdf can be created using latex or converting a word document to pdf or any other method you prefer, as long as it is organized and easy to read. The submission template with latex and code for this homework can be found here: https://classroom.github.com/a/mCG_ijF3.



Info: Szeliski refers to the second edition of *Computer Vision: Algorithms and Applications*, which can be found here.

1 Part 1 Questions (30%)

Instructions

- Three graded questions.
- Write code where appropriate.
- Feel free to include images or equations.
- **We do NOT expect you to fill up each page with your answer.** Some answers will only be a few sentences long, and that is okay.

1.1 Q1

The Harris Corner Detector is commonly used in computer vision algorithms to find interest points from which to extract stable features for image matching.

How do the eigenvalues of the 'M' matrix change if we apply a low-pass filter to the image? What if we apply a high-pass filter? Describe qualitatively.

Answer: If we apply a low-pass filter to the image, the eigenvalues of the M matrix will **decrease**. This is because the low-pass filter will smooth the image, which will reduce the gradient magnitudes and thus the eigenvalues of the M matrix. The opposite is true for a high-pass filter, where the eigenvalues of the M matrix will **increase** because the high-pass filter will keep the high frequencies in the image, which will increase the gradient magnitudes and thus the eigenvalues of the M matrix.

1.2 Q2

Given an interest point location, the SIFT algorithm converts a 16×16 patch around the interest point into a 128×1 feature descriptor of the gradient magnitudes and orientations therein. Write pseudocode *with matrix/array indices* for these steps.

Notes Do this for just one interest point at one scale; ignore the overall interest point orientation; ignore the Gaussian weighting; ignore all normalization post-processing; ignore image boundaries; ignore

sub-pixel interpolation and just pick an arbitrary center within the 16×16 for your feature descriptor. Please just explain in pseudocode how to go from the 16×16 patch to the 128×1 vector.

```
# You can assume access to the image, x and y gradients, and their magnitudes/
# orientations.

image = imread('example.jpg')
grad_x = filter(image, 'sobelX')
grad_y = filter(image, 'sobelY')
grad_mag = sqrt( grad_x.^2 + grad_y.^2 )
grad_ori = atan2( grad_y, grad_x )

# Takes in a interest point x,y location and returns a feature descriptor
def SIFTdescriptor(x, y)
    descriptor = zeros(128,1)

    # loop through a 4x4 grid of 4x4 patches
    for i in range(-2, 2):
        for j in range(-2, 2):
            # i,j represents the 4x4 patch index
            patch_descriptor = zeros(8,1)
            for m in range(4):
                for n in range(4):
                    # m,n represents the 4x4 pixel index
                    # get the gradient orientation and magnitude
                    grad_magnitude = grad_mag[x+i*4+m, y+j*4+n]
                    grad_orientation = grad_ori[x+i*4+m, y+j*4+n]
                    # convert the orientation to the 8-bin histogram index
                    bin_index = floor(grad_orientation // (pi/4))
                    # add the magnitude to the corresponding bin
                    patch_descriptor[bin_index] += grad_magnitude
            # add the 8-bin histogram to the descriptor
            descriptor.append(patch_descriptor)
    return descriptor
```

1.3 Q3

Distance metrics for feature matching.

- Explain the differences between the geometric interpretations of the Euclidean distance and the cosine similarity metrics. What does this tell us about when each should be used?
- Given a distance metric, what is a good method for feature descriptor matching and why?

Answer: Difference between Euclidean distance and cosine similarity: The Euclidean distance measures the distance between two points in a space, while the cosine similarity measures the angle between two vectors. The Euclidean distance is sensitive to the magnitude of the vectors, while the cosine similarity is not. This means that the Euclidean distance will be large if the vectors are far apart in space and cosine similarity will be large if the direction of the vectors is vastly different. This tells us that the Euclidean distance should be used when the magnitude of the vectors is important, while the cosine similarity should be used when the direction of the vectors is important.

Good method for feature descriptor matching: